

R Exposure 4

Unit 1: Strings

Chuck Lanfear

Feb 28, 2021

Updated: Feb 25, 2021



Data for this Unit

We'll use data on food safety inspections in King County from data.kingcounty.gov.

Note these data are *slightly large*. You may want to save them and load them from a *local directory*.

```
library(tidyverse)
restaurants <-
  read_csv("https://clanfear.github.io/CSS508/Lectures/Week8/restaurants.csv",
           col_types = "cccccccnccicccciD")
```

I recommend specifying the column types so they read in correctly.

```
glimpse(restaurants)
```

```
## Rows: 258,630
## Columns: 23
## $ Name          <chr> "@ THE SHACK, LLC ", "10 MERCE...
## $ Program_Identifier <chr> "SHACK COFFEE", "10 MERCER RES...
## $ Inspection_Date  <chr> NA, "01/24/2017", "01/24/2017"...
## $ Description      <chr> "Seating 0-12 - Risk Category ...
## $ Address          <chr> "2920 SW AVALON WAY", "10 MERC...
## $ City             <chr> "Seattle", "Seattle", "Seattle...
## $ Zip_Code         <chr> "98126", "98109", "98109", "98...
## $ Phone            <chr> "(206) 938-5665", NA, NA, NA, ...
## $ Longitude        <dbl> -122, -122, -122, -122, -122, ...
## $ Latitude         <dbl> 47.6, 47.6, 47.6, 47.6, 47.6, ...
## $ Inspection_Business_Name <chr> NA, "10 MERCER RESTAURANT", "1...
## $ Inspection_Type   <chr> NA, "Routine Inspection/Field ...
## $ Inspection_Score  <int> NA, 10, 10, 10, 15, 15, 15, 0,...
## $ Inspection_Result <chr> NA, "Unsatisfactory", "Unsatis...
## $ Inspection_Closed_Business <chr> NA, "false", "false", "false",...
## $ Violation_Type    <chr> NA, "blue", "blue", "red", "bl...
## $ Violation_Description <chr> NA, "4300 - Non-food contact s...
## $ Violation_Points  <int> 0, 3, 2, 5, 5, 5, 5, 0, 5, 10,...
## $ Business_ID       <chr> "PR0048053", "PR0049572", "PR0...
## $ Inspection_Serial_Num <chr> NA, "DAHSIBSJT", "DAHSIBSJT", ...
## $ Violation_Record_ID <chr> NA, "IV43WZVLN", "IVCQ1ZIV0", ...
## $ Grade            <int> NA, 2, 2, 2, 2, 2, 2, 2, 2, 2,...
## $ Date             <date> NA, 2017-01-24, 2017-01-24, 2...
```

Strings

A general programming term for a unit of character data is a **string**, which is defined as *a sequence of characters*. In R the terms "strings" and "character data" are mostly interchangeable.

In other languages, "string" often also refers to a *sequence* of numeric information, such as binary strings (e.g. "01110000 01101111 01101111 01110000"). We rarely use these in R.

Note that these are *sequences* of numbers rather than single numbers, and thus *strings*.

One thing that separates a string from a number is that the leading zeroes are meaningful: `01 != 1`

String Basics

nchar()

`nchar()` gets the number of characters in a string. How many characters are in the ZIP codes?

```
restaurants %>%  
  mutate(ZIP_length = nchar(Zip_Code)) %>%  
  count(ZIP_length)
```

```
## # A tibble: 2 x 2  
##   ZIP_length      n  
##   <int>    <int>  
## 1         5 258629  
## 2        10      1
```

substr()

`substr()` extracts *sub strings* from a larger string. We can use it to pull out just the first 5 digits of the ZIP code.

```
restaurants <- restaurants %>%  
  mutate(ZIP_5 = substr(Zip_Code, 1, 5))  
restaurants %>% distinct(ZIP_5) %>% head()
```

```
## # A tibble: 6 x 1  
##   ZIP_5  
##   <chr>  
## 1 98126  
## 2 98109  
## 3 98101  
## 4 98032  
## 5 98102  
## 6 98004
```

paste()

We can combine parts of strings together using the `paste()` function, e.g. to make a whole mailing address:

```
restaurants <- restaurants %>%  
  mutate(mailing_address =  
    paste(Address, ", ", City, ", WA ", ZIP_5, sep = ""))  
restaurants %>% distinct(mailing_address) %>% head()
```

```
## # A tibble: 6 x 1  
##   mailing_address  
##   <chr>  
## 1 2920 SW AVALON WAY, Seattle, WA 98126  
## 2 10 MERCER ST, Seattle, WA 98109  
## 3 1001 FAIRVIEW AVE N Unit 1700A, SEATTLE, WA 98109  
## 4 1225 1ST AVE, SEATTLE, WA 98101  
## 5 18114 E VALLEY HWY, KENT, WA 98032  
## 6 121 11TH AVE E, SEATTLE, WA 98102
```


paste0()

`paste0()` is a shortcut for `paste()` without any separator.

```
paste(1:5, letters[1:5]) # sep is a space by default
```

```
## [1] "1 a" "2 b" "3 c" "4 d" "5 e"
```

```
paste(1:5, letters[1:5], sep = "")
```

```
## [1] "1a" "2b" "3c" "4d" "5e"
```

```
paste0(1:5, letters[1:5])
```

```
## [1] "1a" "2b" "3c" "4d" "5e"
```

paste() Practice

`sep=` controls what happens when doing entry-wise squishing of vectors you give to `paste()`, while `collapse=` controls if/how they go from a vector to a single string.

Here are some examples; make sure you understand how each set of arguments produces its results:

```
paste(letters[1:5], collapse = "!")
paste(1:5, letters[1:5], sep = "+")
paste0(1:5, letters[1:5], collapse = "???)
paste(1:5, "Z", sep = "*")
paste(1:5, "Z", sep = "*", collapse = " ~ ")
```

```
## [1] "a!b!c!d!e"
## [1] "1+a" "2+b" "3+c" "4+d" "5+e"
## [1] "1a???2b???3c???4d???5e"
## [1] "1*Z" "2*Z" "3*Z" "4*Z" "5*Z"
## [1] "1*Z ~ 2*Z ~ 3*Z ~ 4*Z ~ 5*Z"
```

stringr



stringr

`stringr` is yet another R package from the Tidyverse (like `ggplot2`, `dplyr`, `tidyr`, `lubridate`, `readr`).

It provides functions that:

- Replace some basic string functions like `paste()` and `nchar()` in a way that's a bit less touchy with missing values or factors
- Remove whitespace or pad it out
- Perform tasks related to **pattern matching**: Detect, locate, extract, match, replace, split.
 - These functions use **regular expressions** to describe patterns
 - Base R and `stringi` versions for these exist but are harder to use

Conveniently, *most* `stringr` functions begin with "`str_`" to make RStudio auto-complete more useful.

```
library(stringr)
```

stringr Equivalencies

- `str_sub()` is like `substr()` but also lets you put in negative values to count backwards from the end (-1 is the end, -3 is third from end):

```
str_sub("Washington", 1, -3)
```

```
## [1] "Washingt"
```

- `str_c()` ("string combine") is just like `paste()` but where the default is `sep = ""` (like `paste0()`)

```
str_c(letters[1:5], 1:5)
```

```
## [1] "a1" "b2" "c3" "d4" "e5"
```

stringr Equivalencies

- `str_length()` is equivalent to `nchar()`:

```
nchar("weasels")
```

```
## [1] 7
```

```
str_length("weasels")
```

```
## [1] 7
```

Changing Cases

`str_to_upper()`, `str_to_lower()`, `str_to_title()` convert cases, which is often a good idea to do before searching for values:

```
head(unique(restaurants$City))
```

```
## [1] "Seattle" "SEATTLE" "KENT" "BELLEVUE" "KENMORE" "Issaquah"
```

```
restaurants <- restaurants %>%  
  mutate(across(c(Name, Address, City), ~str_to_upper(.)))  
head(unique(restaurants$City))
```

```
## [1] "SEATTLE" "KENT" "BELLEVUE" "KENMORE" "ISSAQUAH" "BURIEN"
```

`str_trim()` Whitespace

Extra leading or trailing whitespace is common in text data:

```
head(unique(restaurants$Name), 4)
```

```
## [1] "@ THE SHACK, LLC "      "10 MERCER RESTAURANT"  
## [3] "100 LB CLAM"           "1000 SPIRITS"
```

Any character column is potentially affected. We can use the `str_trim()` function in `stringr` to clean them up all at once:

```
restaurants <- restaurants %>%  
  mutate(across(where(is.character), ~str_trim(.)))  
head(unique(restaurants$Name), 4)
```

```
## [1] "@ THE SHACK, LLC"      "10 MERCER RESTAURANT"  
## [3] "100 LB CLAM"           "1000 SPIRITS"
```

`across(where(x), ~ y)` applies function `y` to every column for which function `x` returns `TRUE`.

Regular Expressions and Pattern Matching

What are Regular Expressions?

Regular expressions or **regexes** are how we describe patterns we are looking for in text in a way that a computer can understand. We write an **expression**, apply it to a string input, and then can do things with **matches** we find.

- **Literal characters** are defined snippets to search for like `SEA` or `206`
- **Metacharacters** let us be flexible in describing patterns:
 - backslash `\`, caret `^`, dollar sign `$`, period `.`, pipe `|`, question mark `?`, asterisk `*`, plus sign `+`, parentheses `(` and `)`, square brackets `[` and `]`, curly braces `{` and `}`
 - To treat a metacharacter as a literal character, you must **escape** it with two preceding backslashes `\\`, e.g. to match `(206)` including the parentheses, you'd use `\\(206\\)` in your regex

str_detect()

I want to get inspections for coffee shops. I'll say a coffee shop is anything that has "COFFEE", "ESPRESSO", or "ROASTER" in the name. The `regex` for this is `COFFEE|ESPRESSO|ROASTER` because `|` is a metacharacter that means "OR". Use the `str_detect()` function, which returns `TRUE` if it finds what you're looking for and `FALSE` if it doesn't (similar to `grepl()`):

```
coffee <- restaurants %>%  
  filter(str_detect(Name, "COFFEE|ESPRESSO|ROASTER"))  
coffee %>% distinct(Name) %>% head()
```


```
## # A tibble: 6 x 1  
##   Name  
##   <chr>  
## 1 2 SISTERS ESPRESSO  
## 2 701 COFFEE  
## 3 909 COFFEE AND WINE  
## 4 AJ'S ESPRESSO  
## 5 ALKI HOMEFRONT SMOOTHIES & ESPRESSO  
## 6 ALL CITY COFFEE
```

Will My Coffee Kill Me?

Let's take each unique business identifier, keep the most recent inspection score, and look at a histogram of scores:

```
coffee %>% select(Business_ID, Name, Inspection_Score, Date) %>%  
  group_by(Business_ID) %>% filter(Date == max(Date)) %>%  
  distinct(.keep_all=TRUE) %>% ggplot(aes(Inspection_Score)) +  
  geom_histogram(bins=8) + xlab("Most recent inspection score") + ylab("") +  
  ggtitle("Histogram of inspection scores for Seattle coffee shops")
```

Histogram of inspection scores for Seattle coffee shops



str_detect(): Patterns

Let's look for phone numbers whose first three digits are "206" using `str_detect()`.

We will want it to work whether they have parentheses around the beginning or not, but NOT to match "206" occurring elsewhere:

```
area_code_206_pattern <- "^\\(?:206"  
phone_test_examples <- c("2061234567", "(206)1234567",  
                          "(206) 123-4567", "555-206-1234")  
str_detect(phone_test_examples, area_code_206_pattern)
```

```
## [1] TRUE TRUE TRUE FALSE
```

- `^` is a metacharacter meaning "look only at the *beginning* of the string"
- `\\(?:` means look for a left parenthesis (`\\(`), but it's optional (`?`)
- `206` is the literal string to look for after the optional parenthesis

str_view()

`stringr` also has a function called `str_view()` that allows you to see in the viewer pane *exactly* what text is being selected with a regular expression.

```
str_view(phone_test_examples, area_code_206_pattern)
```

This will generate a small web page in the viewer pane (but not in Markdown docs).

Just be careful to not load an entire long vector / variable or it may crash RStudio as it tries to render a massive page!

str_detect()

Perhaps we want to know how many phone numbers aren't in the 206 area code?

```
restaurants %>%  
  mutate(has_206_number =  
    str_detect(Phone, area_code_206_pattern)) %>%  
  count(has_206_number)
```

```
## # A tibble: 3 x 2  
##   has_206_number      n  
##   <lgl>          <int>  
## 1 FALSE          66655  
## 2 TRUE           109099  
## 3 NA             82876
```

`str_detect()` returns `NA` for rows with missing (`NA`) phone numbers--you can't search for text in a missing value.

str_extract()

`str_extract()` extracts substrings that match a regex.

Let's extract the directional part of Seattle of addresses: N, NW, SE, none, etc.

```
direction_pattern <- " (N|NW|NE|S|SW|SE|W|E)( |$)"
direction_examples <- c("2812 THORNDYKE AVE W", "512 NW 65TH ST",
                        "407 CEDAR ST", "15 NICKERSON ST")
str_extract(direction_examples, direction_pattern)
```

```
## [1] " W"      " NW "    NA      NA
```

- The first space will match a space character, then
- `(N|NW|NE|S|SW|SE|W|E)` matches one of the directions in the group
- `(|$)` is a group saying either there is a space after, or it's the end of the address string (`$` means the end of the string)

Where are the Addresses?

```
restaurants %>%  
  distinct(Address) %>%  
  mutate(city_region =  
    str_trim(str_extract(Address, direction_pattern))) %>%  
  count(city_region) %>% arrange(desc(n))
```

```
## # A tibble: 9 x 2  
##   city_region      n  
##   <chr>         <int>  
## 1 NE             2086  
## 2 S              1764  
## 3 <NA>           1745  
## 4 N              879  
## 5 SE             868  
## 6 SW             705  
## 7 E              538  
## 8 NW             438  
## 9 W              235
```

A common operation is to `str_extract()` something with extra spaces and then `str_trim()` them out.

str_replace(): Replacing

Maybe we want to do a street-level analysis of inspections (e.g. compare The Ave to Pike Street). How can we remove building numbers?

```
number_pattern <- "[0-9]*-?[A-Z]? (1/2 )?"
number_examples <-
  c("2812 THORNDYKE AVE W", "1ST AVE", "10A 1ST AVE",
    "10-A 1ST AVE", "5201-B UNIVERSITY WAY NE",
    "7040 1/2 15TH AVE NW")
str_replace(number_examples, number_pattern, replacement = "")
```

```
## [1] "THORNDYKE AVE W"      "1ST AVE"              "1ST AVE"
## [4] "1ST AVE"              "UNIVERSITY WAY NE"    "15TH AVE NW"
```

We can also use the shortcut `str_remove()`:

```
str_remove(number_examples, number_pattern)
```

```
## [1] "THORNDYKE AVE W"      "1ST AVE"              "1ST AVE"
## [4] "1ST AVE"              "UNIVERSITY WAY NE"    "15TH AVE NW"
```

How Does the Building Number regex Work?

Let's break down `"^[0-9]*-?[A-Z]? (1/2)?"`:

- `^[0-9]` means look for a digit between 0 and 9 (`[0-9]`) at the beginning (`^`)
- `*` means potentially match more digits after that
- `-?` means optionally (`?`) match a hyphen (`-`)
- `[A-Z]?` means optionally match (`?`) a letter (`[A-Z]`)
- Then we match a space ()
- `(1/2)?` optionally matches a 1/2 followed by a space since this is apparently a thing with some address numbers

Removing Street Numbers

```
restaurants <- restaurants %>%  
  mutate(street_only = str_remove(Address, number_pattern))  
restaurants %>% distinct(street_only) %>% head(10)
```

```
## # A tibble: 10 x 1  
##   street_only  
##   <chr>  
## 1 SW AVALON WAY  
## 2 MERCER ST  
## 3 FAIRVIEW AVE N UNIT 1700A  
## 4 1ST AVE  
## 5 E VALLEY HWY  
## 6 11TH AVE E  
## 7 112TH AVE NE #125  
## 8 NE BOTHELL WAY  
## 9 NW GILMAN BL C-08  
## 10 NE 20TH ST STE 300
```

How About Units/Suites Too?

Getting rid of unit/suite references is tricky, but a decent attempt would be to drop anything including and after "#", "STE", "SUITE", "SHOP", "UNIT":

```
unit_pattern <- " (#|STE|SUITE|SHOP|UNIT).*$"  
unit_examples <-  
  c("1ST AVE", "RAINIER AVE S #A", "FAUNTLEROY WAY SW STE 108",  
    "4TH AVE #100C", "NW 54TH ST")  
str_remove(unit_examples, unit_pattern)
```

```
## [1] "1ST AVE"          "RAINIER AVE S"    "FAUNTLEROY WAY SW"  
## [4] "4TH AVE"          "NW 54TH ST"
```

How'd the Unit regex Work?

Breaking down " (#|STE|SUITE|SHOP|UNIT).*\$":

- First we match a space
- (#|STE|SUITE|SHOP|UNIT) matches one of those words
- .* matches *any* character (.) after those words, zero or more times (*), until the end of the string (\$)

Removing Units/Suites

```
restaurants <- restaurants %>%  
  mutate(street_only =  
    str_trim(str_remove(street_only, unit_pattern)))  
restaurants %>% distinct(street_only) %>% head(11)
```

```
## # A tibble: 11 x 1  
##   street_only  
##   <chr>  
## 1 SW AVALON WAY  
## 2 MERCER ST  
## 3 FAIRVIEW AVE N  
## 4 1ST AVE  
## 5 E VALLEY HWY  
## 6 11TH AVE E  
## 7 112TH AVE NE  
## 8 NE BOTHELL WAY  
## 9 NW GILMAN BL C-08  
## 10 NE 20TH ST  
## 11 S ORCAS ST
```

Where Does Danger Lurk?

Let's get the number of 45+ point inspections occurring on every street.

```
restaurants %>%  
  filter(Inspection_Score > 45) %>%  
  distinct(Business_ID, Date, Inspection_Score, street_only) %>%  
  count(street_only) %>%  
  arrange(desc(n)) %>%  
  head(n=5)
```

```
## # A tibble: 5 x 2  
##   street_only      n  
##   <chr>        <int>  
## 1 UNIVERSITY WAY NE  108  
## 2 S JACKSON ST      105  
## 3 PACIFIC HWY S      90  
## 4 NE 24TH ST        76  
## 5 RAINIER AVE S      70
```


Other Useful `stringr` Functions

`str_pad(string, width, side, pad)`: Adds "padding" to any string to make it a given minimum width.

`str_subset(string, pattern)`: Returns all elements that contain matches of the pattern.

`str_which(string, pattern)`: Returns numeric indices of elements that match the pattern.

`str_replace_all(string, pattern, replacement)`: Performs multiple replacements simultaneously

End of R4, Unit 1